

MIS: Diagnosing Structural Bottlenecks in Copy-Paste Augmentation for Agricultural Object Detection

Huy Khanh Tran, Yuichiro Nomura, and Hiroshi Mineno
Graduate School of Integrated Science and Technology, Shizuoka University, Japan
tran.huy.khanh.25@shizuoka.ac.jp

Abstract—Copy-paste augmentation achieves strong results for instance segmentation yet yields inconsistent or negative gains for bounding-box detection in agricultural scenes. We propose the Minimum Inference Score (MIS), a difficulty metric derived from a trained baseline detector that correlates with detection failure modes—miss rate and false-positive rate—across objects and images. MIS reveals where a model fails in a dataset: when MIS detects strong or moderate concentration of hard instances in near-zero AP regimes—regardless of whether the structural source is object size, domain composition, or both—copy-paste augmentation cannot provide reliable gains. We demonstrate this as a pre-augmentation diagnostic validated across three agricultural benchmarks, where MIS correctly anticipates failure under structural constraints and consistent gains in their absence.

Index Terms—object detection, copy-paste augmentation, difficulty scoring, agricultural object detection, data augmentation diagnostic

I. INTRODUCTION

Copy-paste augmentation composites cropped labeled instances into training images to increase sampling diversity [1], [2]. While effective for instance segmentation [2], bounding-box detection in agricultural imagery is subject to harsher constraints—sub-resolution instances, multi-domain acquisition, and high object density—frequently producing inconsistent or negative augmentation outcomes and forcing costly multi-seed search without a principled stopping criterion.

We introduce the **Minimum Inference Score (MIS)**, a metric computed from a trained baseline detector that captures model failure patterns on a given dataset. By correlating with miss rate and FP rate, MIS surfaces where difficulty concentrates in a dataset (driven by object size, domain composition or their interaction) and helps determine whether copy-paste can provide useful training signal. On MinneApple [3], GWHHD 2021 [4], and WGISD [5], MIS correctly anticipates failure when hard instances concentrate in structurally constrained regimes and consistent gains when no such constraint is present, providing bidirectional diagnostic validation.

II. METHODOLOGY

A. Scoring Formulation

Let $G(I)$ denote ground-truth boxes in image I . A trained detector produces predictions p with confidence c_p at $\tau_{\text{conf}}=0.001$

to retain near-miss candidates. For each $g \in G(I)$, let $\mathcal{P}(g)=\{p : \text{IoU}(p, g) \geq 0.5\}$. The *Object Difficulty Score* is

$$S_{\text{obj}}(g) = \begin{cases} \min_{p \in \mathcal{P}(g)} [\alpha(1-c_p) + \beta(1-u_p)], & \mathcal{P}(g) \neq \emptyset \\ \alpha + \beta, & \mathcal{P}(g) = \emptyset \end{cases} \quad (1)$$

with $\alpha=\beta=0.5$ and $u_p=\text{IoU}(p, g)$. Unmatched objects receive the maximum score ($\alpha+\beta=1$). The *Image Difficulty Score* is

$$S_{\text{img}}(I) = \frac{1}{|G(I)|} \sum_g S_{\text{obj}}(g) + w_2 \cdot \text{FPrate}(I), \quad (2)$$

where $\text{FPrate}(I)$ counts unmatched predictions per image. This yields $S_{\text{obj}} \in [0, 1]$ for every ground-truth object and S_{img} for every training image, where higher values indicate greater detection difficulty. By construction, S_{img} encodes both axes of detection failure (miss rate and FP rate); w_2 is set by grid search to balance their contributions.

B. Augmentation Protocol

We train a YOLOv8s baseline and partition all training objects by S_{obj} and all training images by S_{img} into equal difficulty tertiles (Low, Medium, High). Four copy-paste conditions are evaluated: **Score High** pastes objects from the highest-difficulty object tertile into the highest-difficulty destination images; **Score Medium** and **Score Low** follow the same logic for their respective tertiles. **Random** selects both destination images and source objects without score-based filtering. This design isolates whether directing augmentation toward specific difficulty regimes affects outcome. Following [1], training data is expanded by 30% via unblended within-image copy-paste: 2–8 objects per image, scale $[0.9, 1.1]$, rotation $\pm 15^\circ$, $3\times$ maximum reuse; all other settings follow the default YOLOv8 protocol.

III. RESULTS

A. MIS Reflects Detection Failure Modes

Pearson correlations between S_{img} and miss rate are $r=0.66, 0.75, 0.63$ on MinneApple, GWHHD, and WGISD; correlations with FP rate are $r=0.63, 0.75, 0.64$. Consistent moderate-to-strong correlations across all three datasets confirm that MIS captures the primary axes of detection failure as intended by its construction.

TABLE I
MINNEAPPLE: S_{obj} CAPTURES SIZE-LEVEL DIFFICULTY.

Category	AP	Mean S_{obj}	Top tertile
Very Small ($\leq 400 \text{ px}^2$)	0.071	0.152	64.4%
Small (400–1024 px^2)	0.270	0.076	33.9%
Medium ($> 1024 \text{ px}^2$)	0.493	0.053	1.7%




Fig. 1. GWHD 2021: per-domain AP vs. median image difficulty (MIS).

B. MIS Identifies Structural Failure Sources

Size (MinneApple). S_{obj} inversely tracks AP across size groups (Spearman $r_s = -0.82$; Table I): 64.4% of the top-difficulty tertile falls in *Very Small* instances ($\leq 400 \text{ px}^2$), where AP is near-zero (0.071), indicating hard objects concentrate below the model’s size recognition threshold.

Domain (GWHD 2021). Per-domain median MIS strongly anticorrelates with domain AP (Pearson $r = -0.90$; Fig. 1): MIS correctly ranks domain difficulty from detector behavior alone, regardless of its underlying cause—varied cameras, growth stages, and geographic acquisition conditions.

C. Application: Pre-Augmentation Diagnostic for Copy-Paste

MinneApple (resolution ceiling). MIS identifies hard objects concentrated below the recognition threshold. All four copy-paste conditions yield negative AP (Table II), confirming that additional sub-resolution instances provide no usable learning signal regardless of selection strategy.

GWHD 2021 (domain difficulty). MIS partitioning reveals near-complete domain separation: low-difficulty images are 66% in ETHZ_1/Rres_1 domain with ($\text{AP} \approx 0.70$), while high-difficulty images are 78% Arvalis variants ($\text{AP} \approx 0.40$ – 0.47). This structure directly predicts per-condition behavior. Score Low sources from high-AP domains, reinforcing already well-learned patterns which result in consistent, low-variance drop (-0.034 , $\sigma = 0.005$). Score High sources from hard Arvalis domains where difficulty is acquisition-based; added instances cannot resolve domain gaps, yielding high instability (-0.006 , $\sigma = 0.040$). Random and Medium sample without domain awareness, producing moderate drops (-0.015 , -0.017) at intermediate variance ($\sigma = 0.023$, 0.021). Outcome variance scales monotonically with sourced-domain difficulty—a pattern that we can use MIS to predicts from partition composition alone, before any augmented run.

TABLE II
COPY-PASTE OUTCOMES (3-SEED MEAN $\pm \sigma$). BOTTLENECK: STRUCTURAL FAILURE SOURCE IDENTIFIED BY MIS.

Dataset	Bottleneck	Base AP	Best Δ	Worst Δ
MinneApple	Size	0.353 ± 0.005	-0.003 (Hi)	-0.008 (Rnd)
GWHD 2021	Domain	0.229 ± 0.013	-0.006^* (Hi)	-0.034 (Lo)
WGISD	—	0.480 ± 0.007	$+0.009$ (Med)	$+0.001$ (Hi)

*Score High on GWHD: $\sigma = 0.040$, high seed variance.

WGISD (positive control). MIS flags no structural bottleneck: unlike MinneApple and GWHD, S_{obj} shows no dominant concentration in near-zero AP size or domain regimes, and baseline AP is non-negligible across all object groups. All conditions yield positive gains (Table II), with Score Medium achieving the best result, confirming copy-paste viability when MIS raises no flag.

Across all three benchmarks, the *direction* of copy-paste outcome is consistent with the MIS diagnostic: all conditions are negative on flagged datasets and positive on the unflagged dataset. This directional consistency, independent of selection strategy, supports a structural rather than selection-driven explanation and directly addresses why copy-paste struggles: it cannot overcome the underlying difficulty MIS identifies.

IV. CONCLUSION

MIS links baseline detector failure behavior to difficulty concentration patterns, enabling a practical pre-augmentation signal that indicates whether copy-paste is likely to yield reliable gains. In our experiments, 24 augmented training runs across MIS-flagged datasets (MinneApple and GWHD, 4 conditions $\times$ 3 seeds each) yielded zero reliable gains; MIS identifies these datasets without any augmented run, directly saving that compute budget. Practitioners should apply MIS after baseline training: when MIS identifies strong or moderate concentration of high-difficulty instances in near-zero AP regimes, copy-paste augmentation is unlikely to yield reliable gains. This finding is validated bidirectionally across three agricultural benchmarks. Future work includes using MIS to guide instance-level selection when structural constraints are absent, and combining MIS with generative methods to produce information-rich training instances for hard regimes.

ACKNOWLEDGMENT

This work was supported by JST FOREST JPMJFR201B, JSPS KAKENHI JP25H01110.

REFERENCES

- [1] M. Kisantal *et al.*, “Augmentation for small object detection,” *arXiv:1902.07296*, 2019.
- [2] G. Ghiasi *et al.*, “Simple copy-paste is a strong data augmentation method for instance segmentation,” in *Proc. CVPR*, 2021, pp. 2918–2928.
- [3] N. Häni, P. Roy, and V. Isler, “MinneApple: A benchmark dataset for apple detection and segmentation,” *IEEE RA-L*, vol. 5, no. 2, pp. 852–858, 2020.
- [4] E. David *et al.*, “Global wheat head detection 2021,” *Plant Phenomics*, 2021.
- [5] T. T. Santos *et al.*, “Grape detection, segmentation, and tracking using deep neural networks,” *Comput. Electron. Agric.*, vol. 170, p. 105247, 2020.