

# MIS: Diagnosing Bottlenecks in Copy-Paste Augmentation for Agricultural Object Detection

Huy Khanh Tran, Yuichiro Nomura, and Hiroshi Mineno  
Graduate School of Integrated Science and Technology, Shizuoka University, Japan  
tran.huy.khanh.25@shizuoka.ac.jp

**Abstract**—Copy-paste augmentation achieves strong results for instance segmentation yet yields inconsistent or negative gains for bounding-box detection in agricultural scenes. We propose the Minimum Inference Score (MIS), a difficulty metric derived from a trained baseline detector that correlates strongly with detection failure modes—miss rate and false-positive rate—across objects and images. MIS partition analysis reveals whether hard instances concentrate in constrained regimes—driven by object size, domain composition, or both. When such concentration is detected, copy-paste fails to provide reliable gains across all tested selection strategies. Furthermore, by examining which difficulty groups dominate each MIS partition, practitioners can anticipate which selection strategies will produce the most harmful outcomes—all before any augmented training run. We validate both findings across three agricultural benchmarks, providing bidirectional validation across constrained and unconstrained datasets.

**Index Terms**—object detection, copy-paste augmentation, difficulty scoring, agricultural object detection, data augmentation diagnostic

## I. INTRODUCTION

Copy-paste augmentation composites cropped instances into training images to increase diversity [1], [2], yet yields inconsistent or negative gains for bounding-box detection in agricultural imagery, forcing costly multi-seed search without a principled stopping criterion. Critically, this failure persists across all tested selection strategies, suggesting the root cause is a dataset-level bottleneck that no sampling policy can resolve.

We introduce the **Minimum Inference Score (MIS)**, a metric computed from a trained baseline detector that captures model failure patterns on a given dataset. MIS partition analysis enables two practical outputs without any augmented run. First, it determines whether copy-paste can provide useful training signal by identifying whether hard instances concentrate in constrained regimes—validated across MinneApple [3], GWHD 2021 [4], and WGISD [5], where MIS correctly predicts copy-paste failure and positive gains in their absence. Second, it predicts which selection strategies risk harmful outcomes: on GWHD 2021 [4], Score Low sourcing from high-AP partitions produces the worst consistent drop, a pattern visible from partition composition.

## II. METHODOLOGY

### A. Scoring Formulation

Let  $G(I)$  denote ground-truth boxes in image  $I$ . A trained detector produces predictions  $p$  with confidence  $c_p$  at  $\tau_{\text{conf}}=0.001$

to retain near-miss candidates. For each  $g \in G(I)$ , let  $\mathcal{P}(g)=\{p : \text{IoU}(p, g) \geq 0.5\}$ . The *Object Difficulty Score* is

$$S_{\text{obj}}(g) = \begin{cases} \min_{p \in \mathcal{P}(g)} [\alpha(1-c_p) + \beta(1-u_p)], & \mathcal{P}(g) \neq \emptyset \\ \alpha + \beta, & \mathcal{P}(g) = \emptyset \end{cases} \quad (1)$$

with  $\alpha=\beta=0.5$  and  $u_p=\text{IoU}(p, g)$ . Unmatched objects receive the maximum score ( $\alpha+\beta=1$ ). The *Image Difficulty Score* is

$$S_{\text{img}}(I) = \frac{1}{|G(I)|} \sum_g S_{\text{obj}}(g) + w_2 \cdot \text{FPrate}(I), \quad (2)$$

where  $\text{FPrate}(I)$  counts unmatched predictions per image. This yields  $S_{\text{obj}} \in [0, 1]$  for every ground-truth object and  $S_{\text{img}}$  for every training image, where higher values indicate greater detection difficulty. By construction,  $S_{\text{img}}$  encodes both axes of detection failure (miss rate and FP rate);  $w_2$  is set by grid search to balance their contributions. MIS scores are model-conditioned and reflect structural properties through detector failure behavior, not directly.

### B. Augmentation Protocol

We train a YOLOv8s baseline and partition all training objects by  $S_{\text{obj}}$  and all training images by  $S_{\text{img}}$  into equal difficulty tertiles (Low, Medium, High). Four copy-paste conditions are evaluated: **Score High** pastes objects from the highest-difficulty object tertile into the highest-difficulty destination images; **Score Medium** and **Score Low** follow the same logic for their respective tertiles. **Random** selects both destination images and source objects without score-based filtering. This design isolates whether targeting specific difficulty regimes affects outcome. Training data is expanded by 30% via unblended within image copy-paste [1]: 2–8 objects per image, scale  $[0.9, 1.1]$ , rotation  $\pm 15^\circ$ ,  $3 \times$  maximum reuse; all other settings follow default YOLOv8 protocol.

## III. RESULTS

### A. MIS Reflects Detection Failure

Pearson correlations between  $S_{\text{img}}$  and miss rate on the validation set are  $r=0.79, 0.81, 0.91$  on MinneApple, GWHD, and WGISD; correlations with false positive rate are  $r=0.79, 0.81, 0.84$ . These strong correlations confirm that images assigned high  $S_{\text{img}}$  are genuinely harder to detect—the model misses more objects and produces more false positives on them—validating MIS as a reliable proxy for detection difficulty.

TABLE I  
MINNEAPPLE:  $S_{\text{obj}}$  CAPTURES SIZE-LEVEL DIFFICULTY.

| Category                               | AP    | Mean $S_{\text{obj}}$ | Top tertile |
|----------------------------------------|-------|-----------------------|-------------|
| Very Small ( $\leq 400 \text{ px}^2$ ) | 0.071 | 0.152                 | 64.4%       |
| Small (400–1024 $\text{px}^2$ )        | 0.270 | 0.076                 | 33.9%       |
| Medium ( $> 1024 \text{ px}^2$ )       | 0.493 | 0.053                 | 1.7%        |




Fig. 1. GWHD 2021: per-domain AP vs. median image difficulty score.

### B. MIS Identifies Dataset Failure Sources

**Size (MinneApple).**  $S_{\text{obj}}$  strongly correlates with object size (Spearman  $r_s = -0.82$ ): 64.4% of the top-difficulty tertile falls in *Very Small* instances ( $\leq 400 \text{ px}^2$ ), where AP is low indicating that the model struggles primarily with small objects.

**Domain (GWHD 2021).** Per-domain median MIS strongly correlates with domain AP (Pearson  $r = -0.90$ ; Fig. 1): MIS assigns higher scores to images from harder acquisition domains, purely from detector behavior, without access to domain labels during scoring.

### C. Application: Copy-Paste Diagnostic and Strategy Risk

**MinneApple (resolution ceiling).** MIS partition analysis shows 64.4% of high-difficulty objects fall in the *Very Small* category ( $\leq 400 \text{ px}^2$ ,  $\text{AP} = 0.071$ ). This signals a resolution ceiling: pasting additional objects provides no useful learning signal because the model already fails on them structurally, not from lack of examples. This is confirmed: all copy-paste conditions yield negative  $\Delta\text{AP}$  (Table II), regardless of selection strategy, with no condition reversing the outcome direction.

**GWHD 2021 (domain difficulty and strategy risk).** MIS partitioning reveals near-complete domain separation: low-difficulty images are 66% ETHZ\_1/Rres\_1 ( $\text{AP} \approx 0.65$ ), while high-difficulty images are 78% Arvalis variants ( $\text{AP} \approx 0.46$ ). Since difficulty comes from the image-level properties of each domain, not from a lack of instances, copy-paste cannot help: it adds more instances but cannot change how each domain looks to the detector. This is confirmed: all conditions yield negative  $\Delta\text{AP}$ . Partition composition further predicts per-strategy risk by revealing which domains each strategy sources from. Score Low draws exclusively from high-AP domains, flooding training with already well-represented patterns—confirmed by the consistent drop ( $-0.034$ ,  $\sigma = 0.005$ ). Score High draws from hard Arvalis domains, introducing high-variance signal that

TABLE II  
COPY-PASTE OUTCOMES (3-SEED MEAN  $\pm \sigma$ ). BOTTLENECK: FAILURE SOURCE IDENTIFIED BY MIS.

| Dataset    | Bottleneck | Base AP           | Best $\Delta$   | Worst $\Delta$ |
|------------|------------|-------------------|-----------------|----------------|
| MinneApple | Size       | $0.353 \pm 0.005$ | $-0.003$ (Hi)   | $-0.008$ (Rnd) |
| GWHD 2021  | Domain     | $0.229 \pm 0.013$ | $-0.006^*$ (Hi) | $-0.034$ (Lo)  |
| WGISD      | —          | $0.480 \pm 0.007$ | $+0.009$ (Med)  | $+0.001$ (Hi)  |

\*Score High on GWHD:  $\sigma = 0.040$ , high seed variance.

the model cannot reliably exploit, yielding the seed instability ( $-0.006$ ,  $\sigma = 0.040$ ). Random and Medium source from mixed partitions, producing moderate drops ( $-0.015$ ,  $-0.017$ ) at variance ( $\sigma = 0.023$ ,  $0.021$ ).

**WGISD (positive control).** MIS raises no flag: difficulty is distributed across size categories without concentration, and no domain separation is present. This signals that copy-paste can introduce useful learning signal. This is confirmed: all conditions yield positive  $\Delta\text{AP}$  (Table II), with Score Medium achieving the best result.

## IV. CONCLUSION

Copy-paste augmentation fails when hard instances concentrate in structurally constrained regimes—below the detector’s resolution range or within acquisition-defined domains that instance-level augmentation cannot bridge. MIS makes these constraints visible before any augmented run: on MinneApple and GWHD 2021, partition analysis correctly predicted failure across all four selection strategies; on WGISD, the absence of such concentration correctly predicted positive gains across all conditions—providing bidirectional validation that MIS functions as a diagnostic, not merely a failure detector. On domain-structured datasets, MIS partition composition delivers a second output: it identifies which selection strategies risk, from partition composition alone, before any augmented run. In our experiments, 24 augmented training runs across MIS-flagged datasets yielded zero reliable gains; MIS identifies these cases without any augmented run, directly saving that compute budget. Future work includes using MIS to guide instance-level selection when structural constraints are absent, and combining MIS with generative methods to produce information-rich instances for hard regimes that copy-paste cannot reach.

## ACKNOWLEDGMENT

This work was supported by JST FOREST JPMJFR201B, JSPS KAKENHI JP25H01110.

## REFERENCES

- [1] M. Kisantal *et al.*, “Augmentation for small object detection,” *arXiv:1902.07296*, 2019.
- [2] G. Ghiasi *et al.*, “Simple copy-paste is a strong data augmentation method for instance segmentation,” in *Proc. CVPR*, 2021, pp. 2918–2928.
- [3] N. Häni, P. Roy, and V. Isler, “MinneApple: A benchmark dataset for apple detection and segmentation,” *IEEE RA-L*, vol. 5, no. 2, pp. 852–858, 2020.
- [4] E. David *et al.*, “Global wheat head detection 2021,” *Plant Phenomics*, 2021.
- [5] T. T. Santos *et al.*, “Grape detection, segmentation, and tracking using deep neural networks,” *Comput. Electron. Agric.*, vol. 170, p. 105247, 2020.